

In [1]:

```
import warnings #Importing few important python libraries
warnings.filterwarnings('ignore')
import pandas as pd
import numpy as np
from sklearn import preprocessing
from numpy import linalg as LA
from sklearn.model_selection import train_test_split
import math
import matplotlib.pyplot as plt
import copy
import math
%matplotlib inline
```

In [2]:

```
df = pd.read_csv('risk_factors_cervical_cancer.csv')
```

In [3]:

```
df.head()
```

Out[3]:

	Age	Number of sexual partners	First sexual intercourse	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	Hormonal Contraceptives
0	18	4.0	15.0	1.0	0.0	0.0	0.0	0.0
1	15	1.0	14.0	1.0	0.0	0.0	0.0	0.0
2	34	1.0	?	1.0	0.0	0.0	0.0	0.0
3	52	5.0	16.0	4.0	1.0	37.0	37.0	1.0
4	46	3.0	21.0	4.0	0.0	0.0	0.0	1.0

5 rows × 36 columns

In [4]:

```
df = pd.read_csv('risk_factors_cervical_cancer.csv',na_values = ["?"])
```

In [5]:

```
df.shape
```

Out[5]:

(858, 36)

In [6]:

```
df.head()
```

Out[6]:

	Age	Number of sexual partners	First sexual intercourse	Num of pregnancies	Smokes	Smokes (years)	Smokes (packs/year)	Hormonal Contraceptives
0	18	4.0	15.0	1.0	0.0	0.0	0.0	0.0
1	15	1.0	14.0	1.0	0.0	0.0	0.0	0.0
2	34	1.0	NaN	1.0	0.0	0.0	0.0	0.0
3	52	5.0	16.0	4.0	1.0	37.0	37.0	1.0
4	46	3.0	21.0	4.0	0.0	0.0	0.0	1.0

5 rows × 36 columns



In [7]:

```
df.dtypes
```

Out[7]:

Age	int64
Number of sexual partners	float64
First sexual intercourse	float64
Num of pregnancies	float64
Smokes	float64
Smokes (years)	float64
Smokes (packs/year)	float64
Hormonal Contraceptives	float64
Hormonal Contraceptives (years)	float64
IUD	float64
IUD (years)	float64
STDs	float64
STDs (number)	float64
STDs:condylomatosis	float64
STDs:cervical condylomatosis	float64
STDs:vaginal condylomatosis	float64
STDs:vulvo-perineal condylomatosis	float64
STDs:syphilis	float64
STDs:pelvic inflammatory disease	float64
STDs:genital herpes	float64
STDs:molluscum contagiosum	float64
STDs:AIDS	float64
STDs:HIV	float64
STDs:Hepatitis B	float64
STDs:HPV	float64
STDs: Number of diagnosis	int64
STDs: Time since first diagnosis	float64
STDs: Time since last diagnosis	float64
Dx:Cancer	int64
Dx:CIN	int64
Dx:HPV	int64
Dx	int64
Hinselmann	int64
Schiller	int64
Citology	int64
Biopsy	int64
dtype:	object

In [8]:

```
# Remove columns with all na values
all_na = df.columns[df.isna().all()]
df.drop(all_na, axis = 1, inplace = True)
```

In [9]:

```
# Separating binary and non binary Data
df_bin = df.loc[:,df.nunique() <=2]
df_nbin = df.loc[:,df.nunique() >2]
```

In [10]:

```
df_bin.shape, df_nbin.shape
```

Out[10]:

```
((858, 24), (858, 12))
```

In [11]:

```
# Check na values in binary  
df_bin.isna().sum()
```

Out[11]:

Smokes	13
Hormonal Contraceptives	108
IUD	117
STDs	105
STDs:condylomatosis	105
STDs:cervical condylomatosis	105
STDs:vaginal condylomatosis	105
STDs:vulvo-perineal condylomatosis	105
STDs:syphilis	105
STDs:pelvic inflammatory disease	105
STDs:genital herpes	105
STDs:molluscum contagiosum	105
STDs:AIDS	105
STDs:HIV	105
STDs:Hepatitis B	105
STDs:HPV	105
Dx:Cancer	0
Dx:CIN	0
Dx:HPV	0
Dx	0
Hinselmann	0
Schiller	0
Citology	0
Biopsy	0
dtype: int64	

In [12]:

```
#Filling na values in binary variables
col_bin = df_bin.columns[df_bin.isna().any()]
for col in col_bin:
    df_bin[col].fillna((df_bin[col].mode()[0]), inplace = True)

df_bin.isna().sum()
```

Out[12]:

```
Smokes                                0
Hormonal Contraceptives              0
IUD                                  0
STDs                                0
STDs:condylomatosis                 0
STDs:cervical condylomatosis        0
STDs:vaginal condylomatosis         0
STDs:vulvo-perineal condylomatosis  0
STDs:syphilis                       0
STDs:pelvic inflammatory disease    0
STDs:genital herpes                 0
STDs:molluscum contagiosum          0
STDs:AIDS                           0
STDs:HIV                            0
STDs:Hepatitis B                    0
STDs:HPV                            0
Dx:Cancer                           0
Dx:CIN                              0
Dx:HPV                              0
Dx                                  0
Hinselmann                          0
Schiller                             0
Citology                             0
Biopsy                              0
dtype: int64
```

In [13]:

```
# Check na values in nonbinary
df_nbin.isna().sum()
```

Out[13]:

```
Age                                0
Number of sexual partners          26
First sexual intercourse            7
Num of pregnancies                 56
Smokes (years)                     13
Smokes (packs/year)                13
Hormonal Contraceptives (years)    108
IUD (years)                        117
STDs (number)                      105
STDs: Number of diagnosis           0
STDs: Time since first diagnosis    787
STDs: Time since last diagnosis    787
dtype: int64
```

In [14]:

```
# Filling na values in non binary variables
col_nbin = df_nbin.columns[df_nbin.isna().any()]
for col in col_nbin:
    df_nbin[col].fillna((df_nbin[col].median()), inplace = True)

df_nbin.isna().sum()
```

Out[14]:

```
Age                                0
Number of sexual partners          0
First sexual intercourse           0
Num of pregnancies                 0
Smokes (years)                     0
Smokes (packs/year)                0
Hormonal Contraceptives (years)    0
IUD (years)                        0
STDs (number)                      0
STDs: Number of diagnosis           0
STDs: Time since first diagnosis    0
STDs: Time since last diagnosis    0
dtype: int64
```

In [15]:

```
# Removing columns with same values in all rows
df_bin = df_bin.loc[:, (df_bin != df_bin.loc[0]).any()]
df_nbin = df_nbin.loc[:, (df_nbin != df_nbin.loc[0]).any()]
```

In [16]:

```
# Removing Duplicate Columns
df_bin = df_bin.T.drop_duplicates().T
df_nbin = df_nbin.T.drop_duplicates().T
```

In [17]:

```
# Description of Non Binary Data
df_nbin.describe()
```

Out[17]:

	Age	Number of sexual partners	First sexual intercourse	Num of pregnancies	Smokes (years)	Smokes (packs/year)	Horn Contracep (y)
count	858.000000	858.000000	858.000000	858.000000	858.000000	858.000000	858.00
mean	26.820513	2.511655	16.995338	2.257576	1.201241	0.446278	2.00
std	8.497948	1.644759	2.791883	1.400981	4.060623	2.210351	3.50
min	13.000000	1.000000	10.000000	0.000000	0.000000	0.000000	0.00
25%	20.000000	2.000000	15.000000	1.000000	0.000000	0.000000	0.00
50%	25.000000	2.000000	17.000000	2.000000	0.000000	0.000000	0.50
75%	32.000000	3.000000	18.000000	3.000000	0.000000	0.000000	2.00
max	84.000000	28.000000	32.000000	11.000000	37.000000	37.000000	30.00

In [18]:

```
# Scaling (Normalization) of Numerical variables
min_max_scaler = preprocessing.MinMaxScaler()
nbin = df_nbin.values
nbin_scaled = min_max_scaler.fit_transform(nbin)
df_nbin_scaled = pd.DataFrame(nbin_scaled, columns = df_nbin.columns)
```

In [19]:

```
# Joining binary and non binary DFs to form cleaned DF
df_cleaned = df_bin.join(df_nbin_scaled)
```

In [20]:

```
X = df_cleaned.drop("Dx:Cancer", axis=1) # X is a dataframe containing all the features
y = df_cleaned["Dx:Cancer"].copy() # Y is a dataframe containing only the target variable

X_test = df_cleaned.drop("Dx:Cancer", axis=1)
y_test = df_cleaned["Dx:Cancer"].copy()
```

In [21]:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=1)
print("X_train.shape", X_train.shape, "y_train.shape", y_train.shape)
print("X_test.shape", X_test.shape, "y_test.shape", y_test.shape)
```

```
X_train.shape (643, 33) y_train.shape (643,)
X_test.shape (215, 33) y_test.shape (215,)
```

In [22]:

```
def sigmoid(z):  
    g = 1/(1+np.exp(-z))  
    return g
```

In [23]:

```
def compute_cost(X, y, theta):  
    # Number of training examples  
    m = len(y)  
  
    # Compute the hypothesis function  
    z = np.dot(X, theta)  
    h = sigmoid(z)  
  
    # Compute the cost function  
    J = (-1 / m) * np.sum(y * np.log(h) + (1 - y) * np.log(1 - h))  
  
    return J
```

In [24]:

```
def gradient(X, y, theta):  
    # Number of training examples  
    m = len(y)  
  
    # Compute the hypothesis function  
    z = np.dot(X, theta)  
    h = sigmoid(z)  
  
    # Compute the gradient  
    #function computes the gradient of the cost function with respect to the parameters.  
    #This is done using the formula below  
    grad = (1 / m) * np.dot(X.T, (h - y))  
  
    return grad
```



In [25]:

```
def gradient_descent(X, y, theta, alpha, num_iters):  
    # Number of training examples  
    m = len(y)  
  
    # List to store the cost function values over iterations  
    J_history = []  
  
    for i in range(num_iters):  
        # Compute the hypothesis function and its error  
        z = np.dot(X, theta)  
        h = 1 / (1 + np.exp(-z))  
        error = h - y  
  
        # Compute the gradient and update the parameters  
        grad = (1 / m) * np.dot(X.T, error)  
        theta -= alpha * grad  
  
        # Compute the cost function and append it to the list  
        J_history.append(compute_cost(X, y, theta))  
  
    return theta, J_history
```

In [26]:

```
np.random.seed(0)  
theta = np.random.rand(X_train.shape[1])  
alpha = 0.01  
num_iters = 1000  
theta, J_history = gradient_descent(X_train, y_train, theta, alpha, num_iters)  
compute_cost(X_train, y_train, theta)
```

Out[26]:

0.3341059499223987

In [27]:

```
def predict(X, theta):  
    # Compute the hypothesis function  
    # By computing the hypothesis function for each example in the input feature matrix  
    # These probabilities can then be rounded to the nearest integer to get the predicted labels  
    z = np.dot(X, theta)  
    h = sigmoid(z)  
  
    # Round to nearest integer to get predicted labels  
    y_pred = np.round(h)  
  
    return y_pred
```

